

Predict the Price of Used Corolla - Bagging and Random Forest

Langtao Chen

Initial: Feb 20, 2017 Update: March 30, 2020

Contents

1.Data	2
2. Data Partitioning	3
3. Bagging and Random Forest	3
3.1. Bagging	3
3.2. Random Forest	4

In this example, we'll use bagging and random forest (RF) to predict the price of used Corolla.

1.Data

Import the data from the csv file.

```
# Clean the environment
rm(list = ls())

# Read data file
df <- read.csv("ToyotaCorolla.csv")
```

```
# Show the head of the dataset
head(df)
```

```
##   Price Age   KM FuelType HP MetColor Automatic   CC Doors Weight
## 1 13500  23 46986   Diesel 90        1          0 2000     3   1165
## 2 13750  23 72937   Diesel 90        1          0 2000     3   1165
## 3 13950  24 41711   Diesel 90        1          0 2000     3   1165
## 4 14950  26 48000   Diesel 90        0          0 2000     3   1165
## 5 13750  30 38500   Diesel 90        0          0 2000     3   1170
## 6 12950  32 61000   Diesel 90        0          0 2000     3   1170
```

```
# Show the structure of the dataset
str(df)
```

```
## 'data.frame':   1436 obs. of  10 variables:
## $ Price      : int  13500 13750 13950 14950 13750 12950 16900 18600 21500 12950 ...
## $ Age        : int   23  23  24  26  30  32  27  30  27  23 ...
## $ KM         : int  46986 72937 41711 48000 38500 61000 94612 75889 19700 71138 ...
## $ FuelType   : chr   "Diesel" "Diesel" "Diesel" "Diesel" ...
## $ HP         : int   90  90  90  90  90  90  90  90 192  69 ...
## $ MetColor   : int    1  1  1  0  0  0  1  1  0  0 ...
## $ Automatic  : int    0  0  0  0  0  0  0  0  0  0 ...
## $ CC         : int  2000 2000 2000 2000 2000 2000 2000 2000 1800 1900 ...
## $ Doors      : int    3  3  3  3  3  3  3  3  3  3 ...
## $ Weight     : int  1165 1165 1165 1165 1170 1170 1245 1245 1185 1105 ...
```

```
# Summary statistics
summary(df)
```

```
##      Price      Age      KM      FuelType
## Min.   : 4350   Min.   : 1.00   Min.   :    1   Length:1436
## 1st Qu.: 8450   1st Qu.:44.00   1st Qu.: 43000   Class :character
## Median : 9900   Median :61.00   Median : 63390   Mode  :character
## Mean    :10731   Mean    :55.95   Mean     : 68533
## 3rd Qu. :11950   3rd Qu.:70.00   3rd Qu.: 87021
## Max.    :32500   Max.    :80.00   Max.    :243000
##      HP      MetColor      Automatic      CC
## Min.   : 69.0   Min.   :0.0000   Min.   :0.00000   Min.   :1300
```

```
## 1st Qu.: 90.0    1st Qu.:0.0000    1st Qu.:0.00000    1st Qu.:1400
## Median :110.0    Median :1.0000    Median :0.00000    Median :1600
## Mean   :101.5    Mean   :0.6748    Mean   :0.05571    Mean   :1567
## 3rd Qu.:110.0    3rd Qu.:1.0000    3rd Qu.:0.00000    3rd Qu.:1600
## Max.   :192.0    Max.   :1.0000    Max.   :1.00000    Max.   :2000
##      Doors      Weight
## Min.   :2.000    Min.   :1000
## 1st Qu.:3.000    1st Qu.:1040
## Median :4.000    Median :1070
## Mean   :4.033    Mean   :1072
## 3rd Qu.:5.000    3rd Qu.:1085
## Max.   :5.000    Max.   :1615
```

From the summary statistics, we found that there is no missing value.

2. Data Partitioning

We use a single 80/20% split.

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
set.seed(1234)
trainIndex <- createDataPartition(df$Price, p = .8, list = FALSE)

train_data <- df[ trainIndex,]
test_data  <- df[-trainIndex,]

nrow(train_data)
```

```
## [1] 1150
```

```
nrow(test_data)
```

```
## [1] 286
```

3. Bagging and Random Forest

We can use the `randomForest()` method in the `randomForest` package to implement the bagging and random forests. Bagging is special case of random forest with `mtry = p`.

3.1. Bagging

```
library(randomForest)

## Warning: package 'randomForest' was built under R version 4.0.4

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

# Fit a bagged tree
bag_corolla <- randomForest(Price~., data = train_data,
                             mtry = 9, importance = TRUE)

bag_corolla

##
## Call:
## randomForest(formula = Price ~ ., data = train_data, mtry = 9,      importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 9
##
##              Mean of squared residuals: 1168995
##              % Var explained: 91.22
```

mtry=9 means that all 9 predictors are considered for each split of the tree. Thus, this is a bagged regression tree.

Test the performance of the bagged tree on the test dataset.

```
library(caret)

bag_yhat <- predict(bag_corolla, newdata = test_data)

postResample(bag_yhat, test_data$Price)
```

```
##           RMSE      Rsquared      MAE
## 1158.4105028    0.8937825  827.5579611
```

3.2. Random Forest

To fit a random forest, we can set a different value for the mtry parameter. The default value of mtry is $p/3$ in the randomForest() method. Here let's set mtry = 5.

```
# Fit a random forest
rf_corolla <- randomForest(Price~., data = train_data,
                           mtry = 5, importance = TRUE)

rf_corolla
```

```
##
## Call:
##  randomForest(formula = Price ~ ., data = train_data, mtry = 5,      importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 5
##
##              Mean of squared residuals: 1151938
##              % Var explained: 91.34
```

Test the performance of the random forest on the test dataset.

```
rf_yhat <- predict(rf_corolla, newdata = test_data)


postResample(rf_yhat, test_data$Price)
```

```
##          RMSE      Rsquared      MAE
## 1122.1206570    0.9001429  803.3361902
```

We can see that the random forest has a better prediction performance than the bagged regression tree.

```
# Plot the importance
varImpPlot(rf_corolla)
```

rf_corolla



We can see that age of the used corolla is the most important feature. If age is excluded from the model, the increase in MSE and node purity will be very large. The three most important features are age, weight, and KM.

Let's further tune the parameter `mtry` by using a repeated 10-fold cross-validation. Here, we use the `train()` method in the `caret` package.

```
tuneGrid <- data.frame(mtry = 1:9)
```

```
tuneGrid
```

```
##   mtry
## 1    1
## 2    2
## 3    3
## 4    4
## 5    5
## 6    6
## 7    7
## 8    8
## 9    9
```

```
library(lubridate)
```

```
## Warning: package 'lubridate' was built under R version 4.0.4
```

```

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union

control <- trainControl(method = 'repeatedcv',
                        number = 10,
                        repeats = 3)

set.seed(123)

# print out system time before training
start_t <- Sys.time()
cat("",cat("Training started at:",format(start_t, "%a %b %d %X %Y")))

## Training started at: Tue Apr 06 3:50:54 PM 2021

rf_tuned <- train(Price ~ ., data = train_data,
                 method = 'rf',
                 trControl = control,
                 tuneGrid = tuneGrid)

# print out system time after training
finish_t <- Sys.time()
cat("",cat("Training finished at:",format(finish_t, "%a %b %d %X %Y")))

## Training finished at: Tue Apr 06 3:56:01 PM 2021

cat("The training process finished in",difftime(finish_t,start_t,units="mins"), "minutes")

## The training process finished in 5.122552 minutes


print(rf_tuned)

## Random Forest
##
## 1150 samples
##    9 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 1035, 1035, 1035, 1036, 1035, 1034, ...
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  1     1850.990  0.8480542  1398.2980
##  2     1187.009  0.8997856   892.0086
##  3     1098.338  0.9094517   831.9544
##  4     1078.569  0.9118771   824.1121
##  5     1073.734  0.9124241   824.8960

```

```
## 6      1069.176  0.9130877  824.5403
## 7      1068.019  0.9131663  825.5026
## 8      1070.270  0.9127050  828.1985
## 9      1071.552  0.9125353  828.7931
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 7.
```

```
plot(rf_tuned)
```



The cross-validation shows that $mtry = 7$ leads to the best model. However, the difference between $mtry = 5$ and $mtry = 7$ is not very large. Our initial choice $mtry = 5$ is not bad!