

Classification Trees

Langtao Chen

Initial: March 1, 2019 Update: April 6, 2021

Contents

1. Data	2
2. Split Data into Training and Test Sets	3
3. Train A Classification Tree	4
4. Test Performance of the Classification Tree	6
5. Prune the Tree	7

1. Data

As an example, let's use the customer churn dataset to demonstrate how to implement classification tree.

```
df <- read.csv('Telco-Customer-Churn.csv', stringsAsFactors = TRUE)
str(df)
```

```
## 'data.frame':    7043 obs. of  21 variables:
## $ CustomerID      : Factor w/ 7043 levels "0002-ORFBO","0003-MKNFE",...: 5376 3963 2565 5536 6512 6512 ...
## $ Gender          : Factor w/ 2 levels "Female","Male": 1 2 2 2 1 1 2 1 1 2 ...
## $ SeniorCitizen   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Partner         : Factor w/ 2 levels "No","Yes": 2 1 1 1 1 1 1 1 2 1 ...
## $ Dependents      : Factor w/ 2 levels "No","Yes": 1 1 1 1 1 1 2 1 1 2 ...
## $ Tenure          : int  1 34 2 45 2 8 22 10 28 62 ...
## $ PhoneService    : Factor w/ 2 levels "No","Yes": 1 2 2 1 2 2 2 1 2 2 ...
## $ MultipleLines   : Factor w/ 3 levels "No","No phone service",...: 2 1 1 2 1 3 3 2 3 1 ...
## $ InternetService : Factor w/ 3 levels "DSL","Fiber optic",...: 1 1 1 1 2 2 2 1 2 1 ...
## $ OnlineSecurity  : Factor w/ 3 levels "No","No internet service",...: 1 3 3 3 1 1 1 3 1 3 ...
## $ OnlineBackup    : Factor w/ 3 levels "No","No internet service",...: 3 1 3 1 1 1 3 1 1 3 ...
## $ DeviceProtection: Factor w/ 3 levels "No","No internet service",...: 1 3 1 3 1 3 1 1 3 1 ...
## $ TechSupport     : Factor w/ 3 levels "No","No internet service",...: 1 1 1 3 1 1 1 1 3 1 ...
## $ StreamingTV     : Factor w/ 3 levels "No","No internet service",...: 1 1 1 1 1 3 3 1 3 1 ...
## $ StreamingMovies : Factor w/ 3 levels "No","No internet service",...: 1 1 1 1 1 3 1 1 3 1 ...
## $ Contract       : Factor w/ 3 levels "Month-to-month",...: 1 2 1 2 1 1 1 1 1 2 ...
## $ PaperlessBilling: Factor w/ 2 levels "No","Yes": 2 1 2 1 2 2 2 1 2 1 ...
## $ PaymentMethod   : Factor w/ 4 levels "Bank transfer (automatic)",...: 3 4 4 1 3 3 2 4 3 1 ...
## $ MonthlyCharges  : num  29.9 57 53.9 42.3 70.7 ...
## $ TotalCharges    : num  29.9 1889.5 108.2 1840.8 151.7 ...
## $ Churn           : Factor w/ 2 levels "No","Yes": 1 1 2 1 2 2 1 1 2 1 ...
```

The response if the Churn column, which is a binary variable.

```
summary(df)
```

```
##      CustomerID      Gender      SeniorCitizen      Partner      Dependents
## 0002-ORFBO:    1   Female:3488      Min.   :0.0000      No :3641      No :4933
## 0003-MKNFE:    1    Male  :3555      1st Qu.:0.0000      Yes:3402      Yes:2110
## 0004-TLHLJ:    1                                     Median :0.0000
## 0011-IGKFF:    1                                     Mean   :0.1621
## 0013-EXCHZ:    1                                     3rd Qu.:0.0000
## 0013-MHZWF:    1                                     Max.   :1.0000
## (Other)       :7037
##      Tenure      PhoneService      MultipleLines      InternetService
## Min.   : 0.00      No : 682      No           :3390      DSL           :2421
## 1st Qu.: 9.00      Yes:6361      No phone service: 682      Fiber optic:3096
## Median :29.00                                     Yes           :2971      No           :1526
## Mean   :32.37
## 3rd Qu.:55.00
## Max.   :72.00
##
##      OnlineSecurity      OnlineBackup
## No           :3498      No           :3088
```

```

## No internet service:1526 No internet service:1526
## Yes :2019 Yes :2429
##
##
##
##
## DeviceProtection TechSupport
## No :3095 No :3473
## No internet service:1526 No internet service:1526
## Yes :2422 Yes :2044
##
##
##
## StreamingTV StreamingMovies Contract
## No :2810 No :2785 Month-to-month:3875
## No internet service:1526 No internet service:1526 One year :1473
## Yes :2707 Yes :2732 Two year :1695
##
##
##
## PaperlessBilling PaymentMethod MonthlyCharges
## No :2872 Bank transfer (automatic):1544 Min. : 18.25
## Yes:4171 Credit card (automatic) :1522 1st Qu.: 35.50
## Electronic check :2365 Median : 70.35
## Mailed check :1612 Mean : 64.76
## 3rd Qu.: 89.85
## Max. :118.75
##
## TotalCharges Churn
## Min. : 18.8 No :5174
## 1st Qu.: 401.4 Yes:1869
## Median :1397.5
## Mean :2283.3
## 3rd Qu.:3794.7
## Max. :8684.8
## NA's :11

```

Since customer ID is a unique identifier for each customer, let's remove it. There are 11 missing values in TotalCharges column. We remove all missing values from the dataset.

```

df$CustomerID <- NULL
df <- na.omit(df)

```

2. Split Data into Training and Test Sets

Let's split the data into training set (50%) and test set (50%).

```

set.seed(123)
train <- sample(1:nrow(df), nrow(df)*0.5)

```

```
train_df <- df[train,]
test_df <- df[-train,]

# Num of observations in training set
nrow(train_df)
```

```
## [1] 3516
```

```
# Num of observations in test set
nrow(test_df)
```

```
## [1] 3516
```

3. Train A Classification Tree

We use the `tree()` method in the `tree` package to fit a classification tree to the training data. The syntax is very similar as the regression tree fit.

The `tree()` method will fit a regression tree if the response variable is continuous, and a classification tree if the response variable is categorical.

```
library(tree)
```

```
## Warning: package 'tree' was built under R version 4.0.4
```


```
# Fit a decision tree
tree_churn <- tree(Churn ~ ., data = train_df)

# Summary of the decision tree
summary(tree_churn)
```

```
##
## Classification tree:
## tree(formula = Churn ~ ., data = train_df)
## Variables actually used in tree construction:
## [1] "Contract" "InternetService" "Tenure"
## Number of terminal nodes: 6
## Residual mean deviance: 0.8835 = 3101 / 3510
## Misclassification error rate: 0.2162 = 760 / 3516
```

From the summary, we notice that only three variables (i.e., `Contract`, `InternetService`, and `Tenure`) are used to construct the tree.


```
# Plot the decision tree
plot(tree_churn)
text(tree_churn, cex = 0.75, col = 'red')
```



From the above, we can see that factor labels have been shortened. To show the full names of factor levels, set the pretty parameter as FALSE.

```
# Plot the decision tree
plot(tree_churn)

text(tree_churn, cex = 0.75,
      col = 'red', pretty = FALSE)
```



The classification tree predicts a customer will churn her service if she has a month-to-month contract, fiber-optic internet service, and tenure < 13.5.

4. Test Performance of the Classification Tree

As we log transform the response variable, the predicted value of the response needs to be transformed back to the original scale.

```
yhat <- predict(tree_churn, newdata = test_df, type = 'class')
yobs <- test_df$Churn
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
confusionMatrix(yhat, yobs, positive = 'Yes')
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction   No  Yes
```

```
##      No  2461  594
##      Yes  128  333
##
##              Accuracy : 0.7947
##              95% CI : (0.7809, 0.8079)
##      No Information Rate : 0.7363
##      P-Value [Acc > NIR] : 4.719e-16
##
##              Kappa : 0.3694
##
##      McNemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.35922
##              Specificity : 0.95056
##      Pos Pred Value : 0.72234
##      Neg Pred Value : 0.80556
##      Prevalence : 0.26365
##      Detection Rate : 0.09471
##      Detection Prevalence : 0.13111
##      Balanced Accuracy : 0.65489
##
##      'Positive' Class : Yes
##
```

5. Prune the Tree

From the plot of the classification tree shown in section 4, we notice that the tree has unnecessarily complicated structure:

- If the customer has a one year or two year contract, we can simply predict that the customer will not churn her service. So there is no need to further evaluate if the contract is one year or not. That is to say, the 1st and 2nd terminal nodes can be combined into a single one.
- If the customer has a month-to-month contract and a DSL or no internet service, we don't need to further evaluate tenure. The customer churn can be classified as No. That is to say, the 3rd and 4th terminal nodes in the tree plot in section 4 can be combined into a single one.

Thus, the above decision tree can be pruned.

Generally speaking, an unpruned tree may overfit the data (low bias, high variance). By pruning the tree, we may improve its performance.


Now, we check whether pruning the tree might lead to improved interpretation and some times better performance.

```
cv_churn <- cv.tree(tree_churn, FUN=prune.misclass)
cv_churn
```

```
## $size
## [1] 6 4 1
##
## $dev
## [1] 798 798 931
```

```
##
## $k
## [1] -Inf 0.00000 60.66667
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune" "tree.sequence"
```

```
plot(cv_churn$size, cv_churn$dev, type='b',
     xlab = 'Tree Size', ylab = 'Deviance')
```

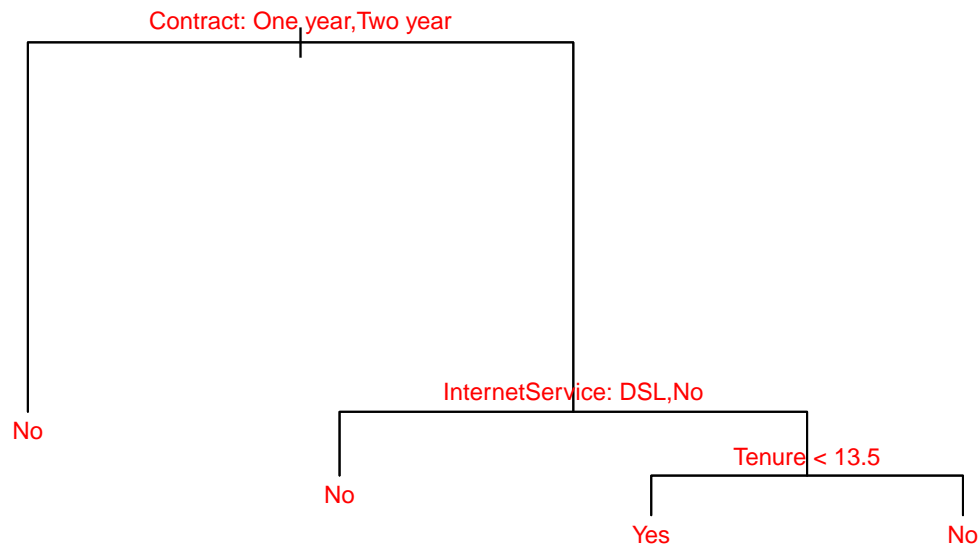


We find the misclassification would be the same if we prune the tree to keep 4 terminal nodes. We can apply the `prune.misclass()` function in to prune the tree.

```
tree_churn_pruned <- prune.misclass(tree_churn, best = 4)

# Plot the pruned tree
plot(tree_churn_pruned)

text(tree_churn_pruned, cex = 0.75,
     col = 'red', pretty = FALSE)
```

We can test the performance of the pruned tree. As shown below, the pruned tree has the same performance as the unpruned tree. However, the pruned tree has a better interpretation as a smaller set of decisions rules are generated from the pruned tree. In addition, the pruned tree tends to have lower variance.

```

yhat2 <- predict(tree_churn_pruned, newdata = test_df, type = 'class')

library(caret)
confusionMatrix(yhat2, yobs, positive = 'Yes')

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction   No  Yes
##           No 2461 594
##           Yes 128 333
##
##           Accuracy : 0.7947
##           95% CI : (0.7809, 0.8079)
##           No Information Rate : 0.7363
##           P-Value [Acc > NIR] : 4.719e-16
##
##           Kappa : 0.3694
##
##           McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.35922

```

```
##          Specificity : 0.95056
##      Pos Pred Value : 0.72234
##      Neg Pred Value : 0.80556
##          Prevalence : 0.26365
##      Detection Rate : 0.09471
## Detection Prevalence : 0.13111
##      Balanced Accuracy : 0.65489
##
##      'Positive' Class : Yes
##
```